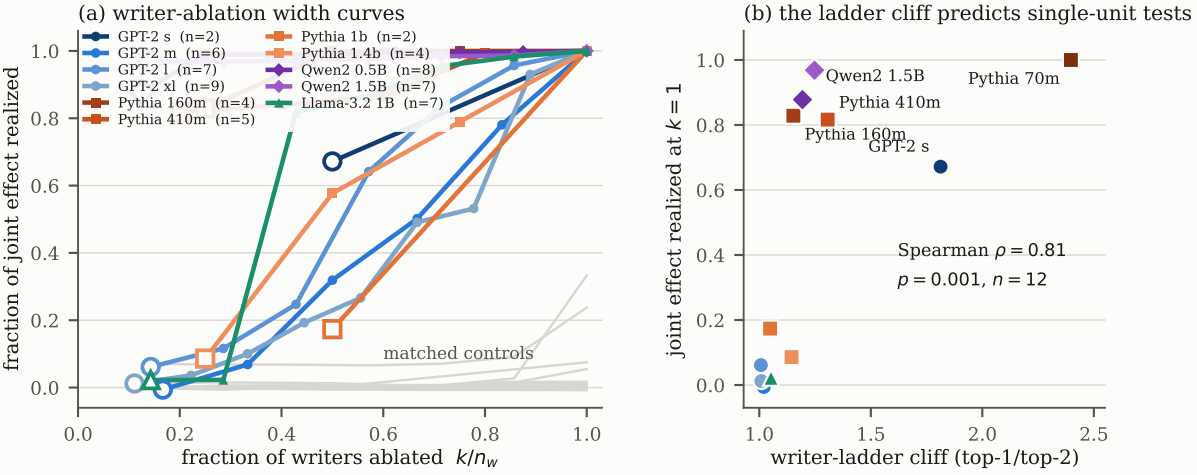
(a) writer-ablation width curves
GPT-2 s (n=2)
GPT-2 m (n=6)
GPT-2 l (n=7)
GPT-2 xl (n=9)
Pythia 160m (n=4)
Pythia 410m (n=5)
Pythia 1b (n=2)
Pythia 1.4b (n=4)
Qwen2 0.5B (n=8)
Qwen2 1.5B (n=7)
Llama-3.2 1B (n=7)
matched controls
fraction of joint effect realized
fraction of writers ablated k/n_w
0.0
0.2
0.4
0.6
0.8
1.0
(b) the ladder cliff predicts single-unit tests
Qwen2 1.5B
Pythia 70m
Pythia 410m
Pythia 160m
GPT-2 s
Spearman $\rho = 0.81$
$p = 0.001$, $n = 12$
joint effect realized at $k = 1$
writer-ladder cliff (top-1/top-2)
1.0
1.5
2.0
2.5